

Lightweight Stacked Ensemble for AI-Enabled 5G Network Intrusion Detection: A Latency–Accuracy Pareto Analysis

1st Van-Quynh Trinh

Information Technology Faculty

Posts and Telecommunications Institute of Technology

Ho Chi Minh City, Vietnam

tvquynh@ptit.edu.vn / ORCID 0009-0006-0514-6123

2nd Trong-Thua Huynh

Information Technology Faculty

Posts and Telecommunications Institute of Technology

Ho Chi Minh City, Vietnam

huynhtrongthua@ptit.edu.vn / ORCID 0000-0003-3934-1067

3rd De-Thu Huynh

Faculty of Information Technology

Saigon International University

Ho Chi Minh City, Vietnam

ORCID 0000-0002-1227-0281

4th Ngoc-Hieu Le

Information Technology Faculty

Posts and Telecommunications Institute of Technology

Ho Chi Minh City, Vietnam

Abstract—The 5G access plane multiplies the attack surface of mobile networks while demanding sub-millisecond decision budgets at the edge. Recent deep models deliver state-of-the-art accuracy on the 5G-NIDD benchmark, but their inference cost and parameter footprint are not always compatible with in-line classification on a commodity CPU. We benchmark four lightweight classifiers on a 60-core CPU host: LightGBM, XGBoost, a two-layer multi-layer perceptron (MLP), and a *heterogeneous shallow stacked ensemble* that combines all three through out-of-fold (OOF) cross-validated probability stacking and a logistic-regression meta-learner. On the 5G-NIDD multi-class task we report classification quality under a stratified random split and a cross-base-station split that stresses generalisation across two physically separated radio cells, and characterise the latency–accuracy frontier across batch sizes spanning $\{1, 16, 64, 256, 1024, 4096\}$. Three findings emerge. First, on the in-distribution random split all four models exceed macro-F1 of 0.998 and the stacked ensemble matches the strongest single base learner. Second, under cross-station shift the MLP becomes the strongest single classifier (macro-F1 0.940, binary-F1 0.957), while the stacked ensemble macro-F1 lies between the gradient-boosted trees and the MLP, indicating that the meta-learner inherits the trees’ shared bias on this dataset. Third, the stacked ensemble exhibits the tightest inter-seed variance ($\sigma = 2 \times 10^{-4}$ in both splits)—a desirable property for predictable deployment even when its mean macro-F1 is not the highest. We translate these findings into a deployment recommendation table that maps operating points (latency budget, accuracy floor, variance tolerance) to the configuration that meets them. The artefact is reviewer-runable.

Index Terms—5G security, network intrusion detection, stacked ensemble, lightweight machine learning, latency-accuracy trade-off, multi-access edge computing.

I. INTRODUCTION

The cellular access plane introduced by 5G is broader, more programmable, and—through control–user-plane separation, network slicing, and multi-access edge computing (MEC)—more exposed than its 4G predecessor. Encryption alone does

not stop volumetric flooding, scanning, or low-and-slow denial of service, and operators routinely call for in-line classifiers that detect malicious flows on the data path with budgets in the tens of microseconds per flow [1], [2]. The public 5G-NIDD benchmark [1], [3] was assembled to drive progress on exactly this regime: it captures over 1.2 million flows of benign and attack traffic from two physically separated base stations in a real 5G test network.

The fast-moving literature on 5G-NIDD [4]–[7] pushes accuracy with deep architectures—transformers, CNN–LSTM hybrids, and deeply stacked neural classifiers—that report macro-F1 values approaching unity on the random split. These gains, however, come with parameter counts and per-flow latencies that complicate deployment on the CPU-only fabric that still dominates user-plane functions in 5G core software stacks. Practitioners therefore ask: *is a heavy deep model the only path to state-of-the-art accuracy on 5G-NIDD, or is a carefully constructed shallow ensemble enough?*

Contributions. This paper makes three contributions.

- 1) We propose a heterogeneous shallow stacked ensemble for 5G-NIDD multi-class intrusion detection (Section IV). The ensemble combines two gradient-boosted tree learners (LightGBM, XGBoost) and a two-layer multi-layer perceptron (MLP) through *out-of-fold* cross-validated probability stacking with a logistic-regression meta-learner—an explicit guard against the “stacker fits the training residuals” leak.
- 2) We report classification quality under a stratified random split and a cross-base-station split BS1→BS2 that re-balances the per-class prior to isolate genuine cross-cell generalisation from prior shift (Section VI).
- 3) We characterise the *latency–accuracy Pareto frontier* for every base learner and the full stack across six

batch sizes on a commodity 60-core CPU host, and emit a deployment recommendation table that maps operating points (latency budget, accuracy floor) to the configuration that meets them.

Reproducibility. Code, configuration files, hyperparameters, the ten project seeds, and the aggregation pipeline are released under the MIT licence at <https://github.com/tvquynh/5g-nidd-stacked-ensemble> (tag `submit-ready-atc-2026`); a single `bash` command reproduces every table and figure on the host described in Section V.

II. RELATED WORK

5G NIDS benchmarks. The 5G-NIDD dataset [1] captures benign and adversarial traffic from two colocated base stations in a real testbed, with eight distinct attack families and a benign class. Earlier intrusion benchmarks such as NSL-KDD [8], UNSW-NB15 [9], CIC-IDS2017 [10], Bot-IoT [11], and Edge-IIoTset [12] are commonly used cross-domain baselines; we focus on 5G-NIDD because of its native 5G radio-plane provenance.

Deep learning on 5G-NIDD. Recent work pushes accuracy with deep architectures. Harshdeep et al. [4] report a transformer-based detector specialised for DDoS variants on 5G-NIDD. Farzaneh et al. [5] apply deep transfer learning, including BiLSTM and Inception base models, to detect DDoS attacks on 5G and beyond networks. Ilias et al. [6] combine convolutional networks with a sparsely gated mixture of experts. Javid et al. [2] propose a lightweight intrusion detection model for DDoS mitigation in 5G and beyond, and Pham et al. [7] introduce ENIDS, a deep-learning ensemble framework for network intrusion detection validated across multiple benchmarks. Our work differs in two respects: (i) we adopt a heterogeneous *shallow* stack (boosted trees + MLP $\rightarrow$ logistic regression) instead of a deep stack of neural blocks, and (ii) we explicitly profile per-sample latency and throughput across batch sizes, an axis under-reported in the deep detector literature.

Stacking. Stacked generalisation [13], [14] combines complementary base models through a meta-learner trained on out-of-fold predictions; in this work the meta-learner is a multi-class logistic regression on the concatenation of three K -dimensional probability vectors. The leakage-free OOF construction is the standard discipline in tabular Kaggle pipelines and has been adopted by recent NIDS benchmarks [15].

III. BACKGROUND AND DATASET

5G-NIDD. We use the author-curated `Encoded.csv` distribution of 5G-NIDD [3], comprising 1 215 890 flows, 89 numeric features (Argus-derived flow statistics augmented with one-hot encodings of TCP flags, protocol, state, and DSCP markings), and nine labels (Benign and eight attack families: ICMPFlood, SYNflood, SYNScan, SlowrateDoS, TCPConnectScan, UDPflood, UDPScan, HTTPflood). The capture spans two base stations (BS1, BS2) on two days and is released under CC BY 4.0.

We add a station marker reconstructed from row position (BTS_1 positions 0–728 315, BTS_2 positions 728 316–1 215 889) and a capture-day flag derived from the published attack-to-day mapping. No class labels were re-scanned or revised.

Evaluation regimes. We adopt two complementary splits. *Random* is a stratified 80/20 split on the multi-class label and represents the in-distribution upper bound. *Cross-station* BS1 $\rightarrow$ BS2 trains on BS1 flows and evaluates on BS2 flows. For each class c , both station-specific subsets are subsampled to $\min(|BS1_c|, |BS2_c|)$ records so the per-station support of class c is matched and station- to-station prior shift inside class c is removed. This matching equalises each class *across stations* but does *not* force a uniform prior across the nine labels themselves; weighted metrics therefore remain sensitive to the residual class-support distribution. Classes for which either $|BS1_c|$ or $|BS2_c|$ is zero are excluded from the rebalanced subset; in 5G-NIDD all nine labels appear in both stations, so no class is silently dropped in practice. The residual gap therefore measures cross-cell generalisation after the class-prior shift has been controlled.

IV. METHOD

A. Heterogeneous Shallow Stacked Ensemble

We instantiate three base learners $\{B_1, B_2, B_3\}$ deliberately chosen for diversity of inductive bias:

- B_1 : LightGBM [16], a leaf-wise gradient-boosted tree ensemble with histogram binning and a softmax multi-class objective.
- B_2 : XGBoost [17] with the histogram tree method, providing a depth-wise gradient-boosting baseline that uses a different regularisation regime than B_1 .
- B_3 : a two-layer multi-layer perceptron ($d \rightarrow 128 \rightarrow 64 \rightarrow K$, ReLU activation, Adam optimiser, early stopping) [18], providing non-tree decision boundaries.

Each B_i exposes a class-probability oracle $B_i : \mathbb{R}^d \rightarrow \Delta^{K-1}$ where $K = 9$ is the label cardinality.

Leakage-free OOF stacking. Given training matrix $X \in \mathbb{R}^{N \times d}$ and labels $y \in \{0, \dots, K-1\}^N$, we partition the training rows into $L = 5$ stratified folds. For each B_i and each fold ℓ , we fit $B_i^{(\ell)}$ on the $L-1$ remaining folds and emit predictions on the held-out fold ℓ . Concatenating across folds yields an out-of-fold (OOF) probability matrix $P_i \in \mathbb{R}^{N \times K}$. We then refit B_i on the full training set and cache its parameters for inference on the test set.

The meta-learner is a multi-class logistic regression $\Phi : \mathbb{R}^{3K} \rightarrow \Delta^{K-1}$ that operates on the horizontal concatenation of P_1, P_2, P_3 . The training target for Φ is y itself; because the columns of P_i are OOF, no row participates in both fitting any $B_i^{(\ell)}$ and supervising Φ . The final predictor on test sample x is

$$\hat{y}(x) = \arg \max_k \Phi([B_1(x); B_2(x); B_3(x)])_k. \quad (1)$$

Computational footprint. The full stack stores three base parameter sets plus a $K \times 3K$ meta weight matrix. Inference

cost is the sum of the three base costs plus a single matrix–vector product and a softmax, dominated by the base learners themselves.

B. Latency–Accuracy Pareto Analysis

We profile inference latency on a held-out test subset under six batch sizes $b \in \{1, 16, 64, 256, 1024, 4096\}$, repeating each measurement ten times after two warm-up passes and reporting the median. We report *per-sample latency* ($\mu\text{s}/\text{flow} = 10^3 \cdot t_{\text{med}}/b$ when t_{med} is the median batch latency in ms) and *throughput* (samples/sec). The shared-host measurement is acknowledged as an upper bound on real-time variance; we mitigate jitter by ten repeats with explicit garbage collection between passes.

Each operating point in the Pareto plot is the pair $(\tilde{\mu}_b, \tilde{F}_1)$ where $\tilde{\mu}_b$ is the median per-sample latency at batch b over the project’s ten seeds and $\tilde{F}_1$ is the median macro-F1 under the same seed distribution.

V. EXPERIMENTAL SETUP

Hardware. All training, inference, and latency profiling run on a single Linux host with two Intel Xeon Gold 6130 CPUs (60 logical cores in total) and 256 GB of RAM, with no GPU. Reported numbers therefore correspond directly to a CPU-only deployment scenario; we do not control CPU affinity or NUMA placement, so the medians should be read as controlled-server estimates rather than hard real-time guarantees.

Software. Python 3.10, scikit-learn 1.7.2, LightGBM 4.6.0, XGBoost 3.2.0, NumPy 2.2.6, pandas 2.3.3. Exact pins are in the released artefact’s `requirements.txt`.

Seeds. We use the ten project seeds $\{42, 123, 456, 789, 1011, 2026, 3141, 4242, 5555, 6789\}$ for all methods and all splits; the data assignment for each split is fully determined by the seed. The cross-station rebalance (Section III) is resampled per seed using a seed-driven NumPy generator, so each seed gives a distinct balanced subset.

Hyperparameters and thread budget. LightGBM: 300 trees, $\eta = 0.05$, 63 leaves, `num_threads=16`, early stopping on 25 rounds. XGBoost: 300 trees, $\eta = 0.05$, depth 8, histogram method, `nthread=16`, early stopping on 25 rounds. MLP: hidden layers (128, 64), Adam, $\alpha = 10^{-4}$, batch size 512, max 40 epochs with early stopping on a 10% internal validation hold-out; sklearn’s MLP delegates threading to the underlying BLAS implementation. Meta-learner: multinomial logistic regression with L2 penalty $C = 1$, the L-BFGS solver, and `n_jobs=-1`. All choices are pinned in `configs/models.yaml`.

VI. RESULTS

A. Classification Quality

Table I reports macro-F1, binary attack recall (binary-TPR), binary-F1, and the binary false-positive rate (binary FPR) across the four model variants and the two splits. Under the stratified random split all models exceed macro-F1 of 0.998, leaving a vanishingly small margin for further improvement; the stacked

ensemble matches the strongest single base (XGBoost) at 0.999. Under the cross-station split BS1→BS2 the four methods diverge. The MLP becomes the strongest single classifier with macro-F1 0.940 ± 0.013 , binary-F1 0.957 ± 0.020 , and binary-TPR 0.919 ± 0.037 , substantially ahead of both gradient-boosted tree variants (macro-F1 0.93, binary-F1 0.887, binary-TPR 0.80). The 12-percentage-point attack-recall advantage of the MLP over the trees is the operationally most consequential gap in this table. The stacked ensemble lands at macro-F1 0.933 ± 0.0002 and binary-F1 0.887—between the trees and the MLP on the mean but with the tightest inter-seed variance by an order of magnitude. The MLP’s binary-F1 advantage is therefore *recall-driven*: it misses fewer attacks at a comparable false-positive rate. The logistic-regression meta-learner inherits the shared tree-bias on the cross-station classes and does not recover the MLP’s transferable representation.

B. Latency–Accuracy Pareto Frontier

Figure 1 shows the Pareto frontier at batch size 256 (representative of MEC line-rate ingest); Figure 2 plots throughput against batch size across two decades. The per-sample latency profile is not flat in the batch size: at very small batches the MLP wins because the forward pass through two dense layers is a fixed-cost matrix multiplication, whereas the tree learners pay the constant overhead of histogram-bin lookups per traversal; at large batches the tree learners (LightGBM in particular) overtake the MLP because their per-sample work scales sub-linearly with the batch under vectorised softmax aggregation while the MLP’s dense kernels saturate. The stacked ensemble adds the sum of its base learners and a negligible $O(K^2)$ softmax over the meta features; the resulting frontier point sits to the right of the bases on the latency axis but at least as high on the accuracy axis, allowing operators to pick by constraint.

C. Per-Class Behaviour Under Cross-Station Shift

Figure 3 and Table III drill into per-class F1 under the cross-station split. A clear trade-off emerges: the MLP excels on the volumetric majority classes (Benign 0.835 vs. 0.69 for the tree-based learners, UDPFlood 0.882 vs. 0.78) but lags on the application-layer DoS classes (HTTPFlood 0.845 vs. 0.97, SlowrateDoS 0.892 vs. 0.96). The gradient-boosted trees show the opposite specialisation. The stacked ensemble follows the tree majority on most classes because two of its three base learners are gradient-boosted trees with similar inductive biases, so the weighted mean of its component probability vectors aligns with the trees and does not recover the MLP’s Benign-class advantage. This explains both the stacked ensemble’s strong agreement with the trees on binary-F1 and its shortfall against the MLP under cross-station shift.

D. Overall Quality Summary

Figure 4 consolidates the macro-F1 numbers in Table I for visual reference.

TABLE I

CLASSIFICATION QUALITY ON THE 5G-NIDD MULTI-CLASS TASK. MEAN $\pm$ STD OVER THE 10 PROJECT SEEDS {42, 123, 456, 789, 1011, 2026, 3141, 4242, 5555, 6789}. *Binary-TPR* (ATTACK RECALL), *Binary-F1*, AND *Binary FPR* ARE COMPUTED AFTER COLLAPSING THE EIGHT ATTACK FAMILIES INTO A SINGLE POSITIVE CLASS AGAINST BENIGN; *Binary FPR* = FP / (FP + TN).

Split	Model	Macro-F1	Binary-TPR	Binary-F1	Binary FPR
Random split	LightGBM	0.998 $\pm$ 0.003	1.000 $\pm$ 0.000	1.000 $\pm$ 0.000	0.000 $\pm$ 0.000
	XGBoost	0.999 $\pm$ 0.000	1.000 $\pm$ 0.000	1.000 $\pm$ 0.000	0.000 $\pm$ 0.000
	MLP	0.998 $\pm$ 0.000	1.000 $\pm$ 0.000	1.000 $\pm$ 0.000	0.000 $\pm$ 0.000
	Stacked ensemble (proposed)	0.999 $\pm$ 0.000	1.000 $\pm$ 0.000	1.000 $\pm$ 0.000	0.000 $\pm$ 0.000
Cross-station (BS1 $\rightarrow$ BS2)	LightGBM	0.929 $\pm$ 0.006	0.797 $\pm$ 0.001	0.887 $\pm$ 0.001	0.003 $\pm$ 0.004
	XGBoost	0.930 $\pm$ 0.003	0.796 $\pm$ 0.001	0.887 $\pm$ 0.000	0.001 $\pm$ 0.000
	MLP	0.940 $\pm$ 0.013	0.919 $\pm$ 0.037	0.957 $\pm$ 0.020	0.002 $\pm$ 0.001
	Stacked ensemble (proposed)	0.933 $\pm$ 0.000	0.797 $\pm$ 0.001	0.887 $\pm$ 0.001	0.000 $\pm$ 0.000




Fig. 1. Latency–accuracy Pareto operating points at batch size 256 on (a) the stratified random split and (b) the cross-station BS1 $\rightarrow$ BS2 split. Per-sample latency on the horizontal axis is logarithmic; each point is the median over the ten project seeds. The random panel is near-saturated; the cross-station panel exposes the operational gap between MLP and the tree-based learners.

TABLE II

INFERENCE LATENCY ON AN INTEL XEON GOLD 6130 CPU SERVER (60 CORES, 256 GB RAM). MEDIAN OVER 10 SEEDS AND 10 TIMING REPETITIONS PER CELL; THROUGHPUT IS SAMPLES PER SECOND, ABBREVIATED $k = 10^3$.

Model	μ s/flow $b = 1$	thr./s $b = 1$	μ s/flow $b = 256$	thr./s $b = 256$	μ s/flow $b = 4096$	thr./s $b = 4096$
LightGBM	3402.2	298	23.6	42.5 k	4.9	205.6 k
XGBoost	31685	32	193.2	5.2 k	15.4	65.4 k
MLP	562.4	1.8 k	4.6	217.9 k	11.9	84.0 k
Stacked ensemble (proposed)	26469	38	197.8	5.1 k	32.3	31.0 k

TABLE III

PER-CLASS F1 ON CROSS-STATION BS1 $\rightarrow$ BS2 (MEDIAN OVER 10 SEEDS). THE MLP EXCELS ON THE VOLUMETRIC CLASSES (BENIGN, UDPFLOOD) WHILE THE GRADIENT-BOOSTED TREES AND THE STACKED ENSEMBLE DOMINATE THE APPLICATION-LAYER DoS CLASSES.

Model	Benign	HTTPFlood	ICMPFlood	SYNFlood	SYNScan	SlowrateDoS	TCPConnectScan	UDPFlood	UDPScan
LightGBM	0.691	0.974	0.995	1.000	0.997	0.957	0.997	0.783	0.996
XGBoost	0.691	0.976	0.999	0.975	0.987	0.959	0.998	0.782	0.998
MLP	0.835	0.845	1.000	0.998	0.997	0.892	0.998	0.882	0.998
Stacked ensemble (proposed)	0.692	0.975	0.999	0.999	0.997	0.957	0.998	0.782	0.998

VII. DISCUSSION: DEPLOYMENT OPERATING POINTS

Reading Tables I and II jointly yields a three-tier deployment recommendation:

- *Latency-critical streaming* (single-flow inference, batch 1): the MLP is the recommended single model. Its forward




Fig. 2. Inference throughput against batch size on the random split. Throughput increases sub-linearly because the dominant per-batch cost plateaus once vectorised kernels saturate.




Fig. 3. Per-class F1 (median over ten seeds) on the cross-station split BS1→BS2. The MLP outperforms the trees on the volumetric classes (Benign, UDP Flood), whereas the trees and the stacked ensemble dominate the application-layer DoS classes.




Fig. 4. Macro-F1 (mean over ten seeds, error bars are 1 std) per model and split.

pass is a fixed-size dense matmul that is an order of magnitude cheaper per flow than the trees, and it is also the strongest cross-station classifier we measured. The stacked ensemble is impractical here because its per-flow cost is the sum of all three bases.

- **Balanced MEC edge** (batch 64–1024): the MLP retains both the latency lead (single-digit microseconds per flow at batch 256) and the accuracy lead under cross-station shift. The stacked ensemble is preferable in this regime

only when *inter-seed reproducibility* is part of the SLO: its macro-F1 standard deviation is two orders of magnitude smaller than the MLP’s (Table I).

- **Accuracy-first analytics** (large batches, ≥ 4096): LightGBM offers the lowest per-sample latency. The choice between LightGBM and the stacked ensemble at large batch is between a small accuracy headroom (stacked) and a $6\times$ latency penalty. For deployments that prioritise binary detection sensitivity under cross-station shift, the MLP remains preferable despite its higher per-batch latency at very large b , since its binary-F1 advantage (0.957 vs. 0.887) translates directly into operational alert quality.

This three-tier mapping is the practical deliverable for operators who must satisfy both an accuracy floor and a latency budget; the artefact publishes the full machine-readable mapping in `results/summary_means.csv` and `results/latency_summary.csv`.

Positioning vs deep detectors. Recent deep detectors on 5G-NIDD [4], [6], [7] pursue peak accuracy with transformer and deeply stacked architectures; we deliberately target the lightweight end of the spectrum where boosted trees and a two-layer MLP can be deployed on a commodity CPU. Direct numerical comparison with those models requires reproducing their published hyperparameters and training pipelines, which lies outside this paper’s page budget. Our position is therefore an honest comparative *latency-accuracy frontier* among practical shallow models rather than a peak-accuracy leaderboard. The empirical finding that the two-layer MLP outperforms both boosted-tree variants and the stacked ensemble on cross-station generalisation is consistent with the deep-learning literature’s preference for representation-learning models when the test distribution differs from the training one.

VIII. THREATS TO VALIDITY

The 5G-NIDD capture covers two base stations and a finite attack catalogue; cross-operator and cross-vendor generalisation remain open. Cross-station numbers are sensitive to the per-class rebalancing scheme; we report a stratified-min rebalance that removes prior shift but does not eliminate concept shift in feature space. We do not include a day-level temporal split because the disjoint attack catalogue between the two capture days in 5G-NIDD entangles temporal drift with attack novelty; a within-day temporal split would require a finer-grained timestamp than is exposed in the public release. Latency measurements were collected on a shared 60-core CPU server with ten timing repetitions per batch size, so the reported medians are controlled-server estimates rather than hard real-time guarantees; absolute latency and throughput will differ on constrained edge-class CPUs and ARM-based MEC platforms, and the relative ordering between models should therefore be re-verified on those targets before deployment. With only ten timing repetitions per batch size we report median and standard deviation but not high-percentile tail latency.

IX. CONCLUSION

We presented an empirical latency–accuracy benchmark of four lightweight classifiers for AI-enabled 5G network intrusion detection on the 5G-NIDD dataset: LightGBM, XGBoost, a two-layer MLP, and a heterogeneous shallow stacked ensemble that combines all three through leakage-free out-of-fold cross-validated probability stacking with a logistic-regression meta-learner. The most actionable finding is that the MLP is the strongest single classifier under cross-station shift in both accuracy (macro-F1 0.940, binary-F1 0.957) and per-flow latency at small batch sizes, while the stacked ensemble matches the trees on the in-distribution random split and offers the lowest inter-seed variance across both splits. The contribution is therefore a calibrated mapping from operating constraints to model choice, rather than a claim of state-of-the-art accuracy. The artefact—code, hyperparameters, ten project seeds, and the aggregation pipeline—is released under MIT with a reviewer-runnable single command to reproduce every number and figure. Future work will profile inference on edge-class CPU and FPGA targets, study calibration- and confidence-weighted meta-learners that could recover the MLP’s binary-detection edge inside a stacked configuration, and extend the cross-station analysis to multi-operator captures as they become available.

ACKNOWLEDGMENTS

This work was supported by the Posts and Telecommunications Institute of Technology (PTIT), Vietnam. The authors used a generative AI assistant for language editing, readability improvement, and code-development support. All AI-assisted outputs were reviewed and verified by the authors, who take full responsibility for the integrity and correctness of the manuscript and accompanying artefact.

REFERENCES

- [1] S. Samarakoon, Y. Siriwardhana, P. Porambage, M. Liyanage, S.-Y. Chang, J. Kim, J. Kim, and M. Ylianttila, “5G-NIDD: A comprehensive network intrusion detection dataset generated over 5G wireless network,” *arXiv preprint arXiv:2212.01298*, 2022.
- [2] I. Javid, S. Khara, J. Frnda *et al.*, “NIDD-enabled lightweight intrusion detection for effective DDoS mitigation in 5G and beyond,” *Scientific Reports*, 2025.
- [3] Y. Siriwardhana, S. Samarakoon, P. Porambage, M. Liyanage, and M. Ylianttila, “5G-NIDD: A real-network intrusion detection dataset generated over a 5G wireless network,” *IEEE Data Descriptions*, vol. 2, 2025.
- [4] K. Harshdeep, S. Konatham, and R. Mathur, “DeepTransIDS: Transformer-based deep learning model for detecting DDoS attacks on 5G NIDD,” *Results in Engineering*, vol. 26, p. 104826, 2025.
- [5] B. Farzaneh, N. Shahriar, A. H. Al Muktadir, M. S. Towhid, and M. S. Khosravani, “DTL-5G: Deep transfer learning-based DDoS attack detection in 5G and beyond networks,” *Computer Communications*, vol. 228, p. 107927, 2024.
- [6] L. Ilias, G. Doukas, V. Lamprou, C. Ntanos, and D. Askounis, “Convolutional neural networks and mixture of experts for intrusion detection in 5G networks and beyond,” *Frontiers in Artificial Intelligence*, 2025, arXiv preprint arXiv:2412.03483, 2024.
- [7] T.-H. Pham *et al.*, “ENIDS: A deep learning-based ensemble framework for network intrusion detection systems,” *IEEE Transactions on Network and Service Management*, 2024.
- [8] S. Hettich and S. D. Bay, “The UCI KDD archive: KDD cup 1999 data,” <http://kdd.ics.uci.edu/databases/kddcup99/>, 1999.
- [9] N. Moustafa and J. Slay, “UNSW-NB15: A comprehensive data set for network intrusion detection systems,” in *Military Communications and Information Systems Conference*, 2015.
- [10] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani, “Toward generating a new intrusion detection dataset and intrusion traffic characterization,” in *Proc. ICISSP*, 2018, pp. 108–116.
- [11] N. Koroniotis, N. Moustafa, E. Sitnikova, and B. Turnbull, “Towards the development of realistic botnet dataset in the Internet of Things for network forensic analytics: Bot-IoT dataset,” *Future Generation Computer Systems*, vol. 100, pp. 779–796, 2019.
- [12] M. A. Ferrag, O. Friha, D. Hamouda, L. Maglaras, and H. Janicke, “Edge-IIoTset: A new comprehensive realistic cyber security dataset of IoT and IIoT applications for centralized and federated learning,” *IEEE Access*, vol. 10, pp. 40 281–40 306, 2022.
- [13] D. H. Wolpert, “Stacked generalization,” *Neural Networks*, vol. 5, no. 2, pp. 241–259, 1992.
- [14] L. Breiman, “Stacked regressions,” *Machine Learning*, vol. 24, no. 1, pp. 49–64, 1996.
- [15] Y. N. Kunang, S. Nurmaini, D. Stiawan, and B. Y. Suprpto, “Attack classification of an intrusion detection system using deep learning and hyperparameter optimization,” *Journal of Information Security and Applications*, 2024.
- [16] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, “LightGBM: A highly efficient gradient boosting decision tree,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [17] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proc. ACM SIGKDD*, 2016, pp. 785–794.
- [18] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg *et al.*, “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.